

210236P\_a03

November 11, 2024

## EN3160 Assignment 3 on Neural Networks

Instructed by Dr. Ranga Rodrigo

Done by Jayakumar W.S. (210236P)

### Contents

<b>EN3160 Assignment 3 on Neural Networks</b>	<b>1</b>
<b>Introduction</b>	<b>2</b>
<b>1 Our own architecture</b>	<b>3</b>
1.1 Single Layer . . . . .	3
1.2 Adding Non-linearity . . . . .	6
1.3 A more efficient implementation . . . . .	9
<b>2 LeNet-5</b>	<b>15</b>
<b>3 Implementing ResNet-18</b>	<b>19</b>
3.1 Finetuning the network . . . . .	19
3.2 Using ResNet-18 as a feature extractor . . . . .	30

## Introduction

This assignment is focused on implementing neural networks for image classification. This is done by using: 1. Our own neural network implementation 2. An implementation of LeNet-5 3. An implementation of ResNet-18

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torchinfo import summary
import matplotlib.pyplot as plt
import gc
```

```
[2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

# 1 Our own architecture

```
[73]: transform = transforms.Compose ([ transforms.ToTensor(), transforms.
      ↪ Normalize((0.5, 0.5, 0.5) , (0.5, 0.5, 0.5))])
batch_size = 32
trainset = torchvision.datasets.CIFAR10(root= './data', train=True,
      ↪ download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
      ↪ shuffle=True, num_workers=2)
testset = torchvision.datasets.CIFAR10(root= './data', train=False,
      ↪ download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
      ↪ shuffle=False, num_workers=2)
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
      ↪ 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

## 1.1 Single Layer

```
[4]: Din = 3*32*32 # Input size (flattened CIFAR=10 image size)
K = 10 # Output size (number of classes in CIFAR=10)
std = 1e-5
# Initialize weights and biases
w = torch.randn(Din, K, device=device, dtype=torch.float, requires_grad=True) *
      ↪ std
b = torch.randn(K, device=device, dtype=torch.float, requires_grad=True)
# Hyperparameters
iterations = 20
lr = 2e-6 # Learning rate
lr_decay = 0.9 # Learning rate decay
reg = 0 # Regularization
loss_history = [ ]
```

```
[5]: for t in range(iterations):
      running_loss = 0.0
      for i, data in enumerate(trainloader, 0):
          # Get inputs and labels
          inputs, labels = data
          Ntr = inputs.shape[0] # Batch size
          x_train = inputs.view(Ntr, -1).to(device) # Flatten input to (Ntr, Din)
          y_train_onehot = nn.functional.one_hot(labels, K).float().to(device) #
          ↪ Convert labels to one-hot

          # Forward pass
          y_pred = x_train.mm(w) + b # Output layer activation
```

```

    # Loss calculation (Mean Squared Error with regularization)
    loss = (1/Ntr) * torch.sum((y_pred - y_train_onehot) ** 2) + reg * _
    ↪ torch.sum(w ** 2)
    running_loss += loss.item()

    # Backpropagation
    dy_pred = (2.0 / Ntr) * (y_pred - y_train_onehot)
    dw = x_train.t().mm(dy_pred) + reg * w
    db = dy_pred.sum(dim=0)

    # Parameter update
    w = w - lr * dw
    b = b - lr * db

    loss_history.append(running_loss / len(trainloader))
    print(f"Epoch {t + 1} / {iterations}, Loss: {running_loss / _
    ↪ len(trainloader)}")

    # Learning rate decay
    lr *= lr_decay

```

```

Epoch 1 / 20, Loss: 11.052271098565841
Epoch 2 / 20, Loss: 9.677958526294047
Epoch 3 / 20, Loss: 9.081760541033608
Epoch 4 / 20, Loss: 8.739522613627896
Epoch 5 / 20, Loss: 8.51057107678912
Epoch 6 / 20, Loss: 8.341256552602873
Epoch 7 / 20, Loss: 8.205948939479015
Epoch 8 / 20, Loss: 8.09767150421289
Epoch 9 / 20, Loss: 8.007261962060813
Epoch 10 / 20, Loss: 7.930841871506879
Epoch 11 / 20, Loss: 7.865830220553788
Epoch 12 / 20, Loss: 7.809729950746815
Epoch 13 / 20, Loss: 7.762098590914286
Epoch 14 / 20, Loss: 7.719234191372237
Epoch 15 / 20, Loss: 7.682436113851771
Epoch 16 / 20, Loss: 7.6493737638111075
Epoch 17 / 20, Loss: 7.620805551513066
Epoch 18 / 20, Loss: 7.594736617570952
Epoch 19 / 20, Loss: 7.573210502997927
Epoch 20 / 20, Loss: 7.552905830449197


```

```

[11]: plt.plot(loss_history)
      plt.xlabel('Iteration')
      plt.ylabel('Loss')
      plt.xticks(range(0, iterations, 2))

```

```
plt.title('Loss history')
plt.show()
```



```
[4]: def calculate_accuracy(dataloader: torch.utils.data.DataLoader, w: torch.
      ↪Tensor, b: torch.Tensor) -> float:
      correct = 0
      total = 0
      with torch.no_grad():
          for data in dataloader:
              inputs, labels = data
              inputs, labels = inputs.to(device), labels.to(device)
              N = inputs.shape[0]
              x = inputs.view(N, -1)
              y = x.mm(w) + b
              predicted = torch.argmax(y, dim=1)
              total += labels.size(0)
              correct += (predicted == labels).sum().item()
      return 100 * correct / total
```

```
[13]: train_accuracy = calculate_accuracy(trainloader, w, b)
      test_accuracy = calculate_accuracy(testloader, w, b)

      print(f"Train accuracy: {train_accuracy:.2f}%")
      print(f"Test accuracy: {test_accuracy:.2f}%")
```

```
Train accuracy: 10.12%
Test accuracy: 10.17%
```

We see above that the performance is extremely poor. This is because the model has no non-linearity. We will add a non-linearity to the model and see if the performance improves. This is added using a hidden layer with sigmoid activation

```
[ ]: del w, b, x_train, y_train_onehot, y_pred, loss, dy_pred, dw, db
      gc.collect()
      if torch.cuda.is_available():
          torch.cuda.empty_cache()
```

## 1.2 Adding Non-linearity

```
[7]: # This implementation is not efficient and is only for educational purposes.␣
      ↪For real-world applications, use PyTorch's built-in functions and classes.␣
      ↪This may fail
      # as memory usage increases with the number of iterations.
```

```
Din = 3*32*32 # Input size (flattened CIFAR=10 image size)
K = 10 # Output size (number of classes in CIFAR=10)
std = 1e-5
# Initialize weights and biases
w1 = torch.randn(Din, 100, device=device, requires_grad=True)
b1 = torch.zeros(100, device=device, requires_grad=True)
w2 = torch.randn(100, K, device=device, requires_grad=True)
b2 = torch.zeros(K, device=device, requires_grad=True)
# Hyperparameters
iterations = 10 # Reduced as memory usage increases
lr = 2e-6 # Learning rate
lr_decay = 0.9 # Learning rate decay
reg = 0 # Regularization
loss_history = [ ]
```

```
[6]: for t in range(iterations):
      running_loss = 0.0
      for i, data in enumerate(trainloader, 0):
          # Get inputs and labels
          inputs, labels = data
          Ntr = inputs.shape[0] # Batch size
          x_train = inputs.view(Ntr, -1).to(device) # Flatten input to (Ntr, Din)
```

```

        y_train_onehot = nn.functional.one_hot(labels, K).float().to(device) #
↪One-hot labels

    # Forward pass
    hidden = x_train.mm(w1) + b1
    hidden_activation = torch.sigmoid(hidden) # Sigmoid activation
    logits = hidden_activation.mm(w2) + b2 # Logits before softmax

    # Compute softmax probabilities
    max_logits = torch.max(logits, dim=1, keepdim=True)[0]
    exp_logits = torch.exp(logits - max_logits)
    probs = exp_logits / torch.sum(exp_logits, dim=1, keepdim=True)

    # Cross-Entropy Loss with L2 regularization
    epsilon = 1e-12 # Small value to prevent log(0)
    log_probs = torch.log(probs + epsilon)
    loss = -torch.sum(y_train_onehot * log_probs) / Ntr
    loss += reg * (torch.sum(w1 ** 2) + torch.sum(w2 ** 2))
    running_loss += loss.item()

    # Backpropagation
    dlogits = (probs - y_train_onehot) / Ntr

    # Gradients for parameters of the second layer
    dw2 = hidden_activation.t().mm(dlogits) + reg * w2
    db2 = dlogits.sum(dim=0)

    # Backpropagate through ReLU activation
    dhidden_activation = dlogits.mm(w2.t())
    dhidden = dhidden_activation * hidden_activation * (1 -
↪hidden_activation) # Derivative of sigmoid

    # Gradients for parameters of the first layer
    dw1 = x_train.t().mm(dhidden) + reg * w1
    db1 = dhidden.sum(dim=0)

    # Parameter updates
    w2 = w2 - lr * dw2
    b2 = b2 - lr * db2
    w1 = w1 - lr * dw1
    b1 = b1 - lr * db1

    loss_history.append(running_loss / len(trainloader))
    print(f"Epoch {t+1} / {iterations}, Loss: {running_loss /
↪len(trainloader)}")

    # Learning rate decay


```

```
lr *= lr_decay
```

```
Epoch 1 / 10, Loss: 10.249201878323742  
Epoch 2 / 10, Loss: 10.233684833890264  
Epoch 3 / 10, Loss: 10.218501103511622  
Epoch 4 / 10, Loss: 10.20577982977576  
Epoch 5 / 10, Loss: 10.194548347861204  
Epoch 6 / 10, Loss: 10.185316179169346  
Epoch 7 / 10, Loss: 10.175520200418191  
Epoch 8 / 10, Loss: 10.168194400143028  
Epoch 9 / 10, Loss: 10.160515865147762  
Epoch 10 / 10, Loss: 10.152873168598743
```

It is observed that the loss values decrease on each iteration

```
[7]: plt.plot(loss_history)  
plt.xlabel('Iteration')  
plt.xticks(range(0, iterations, 2))  
plt.ylabel('Loss')  
plt.title('Training Loss')  
plt.show()
```





```
[8]: def calculate_accuracy(dataloader: torch.utils.data.DataLoader, w1: torch.
    ↪Tensor, b1: torch.Tensor, w2: torch.Tensor, b2: torch.Tensor) -> float:
    correct = 0
    total = 0
    with torch.no_grad():
        for data in dataloader:
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)
            N = inputs.shape[0]
            x = inputs.view(N, -1)
            hidden = torch.sigmoid(x.mm(w1) + b1)
            y = hidden.mm(w2) + b2
            predicted = torch.argmax(y, dim=1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    return 100 * correct / total
```

```
[9]: train_accuracy = calculate_accuracy(trainloader, w1, b1, w2, b2)
    test_accuracy = calculate_accuracy(testloader, w1, b1, w2, b2)

    print(f"Train accuracy: {train_accuracy:.2f}%")
    print(f"Test accuracy: {test_accuracy:.2f}%")
```

Train accuracy: 10.05%  
 Test accuracy: 9.77%

```
[10]: del w1, b1, w2, b2, x_train, y_train_onehot, hidden, hidden_activation, logits,
    ↪probs, log_probs, loss, dlogits, dw2, db2, dhidden_activation, dhidden, dw1,
    ↪db1
    gc.collect()
    if torch.cuda.is_available():
        torch.cuda.empty_cache()
```

### 1.3 A more efficient implementation

```
[13]: Din = 3*32*32 # Input size (flattened CIFAR=10 image size)
    K = 10 # Output size (number of classes in CIFAR=10)
    lr = 1e-3 # Learning rate
    reg = 1e-5 # Regularization strength
```

```
[16]: class NeuralNetwork(nn.Module):
    def __init__(self, Din, H, Dout):
        super(NeuralNetwork, self).__init__()
        self.linear1 = nn.Linear(Din, H)
        self.linear2 = nn.Linear(H, Dout)

        def forward(self, x):
```

```

x = torch.flatten(x, 1)
x = torch.sigmoid(self.linear1(x))
x = self.linear2(x)
return x

```

We will define a function for training and testing the model

```

[17]: def train(model:nn.Module,
            trainloader:torch.utils.data.DataLoader,
            testloader:torch.utils.data.DataLoader,
            iterations:int,
            optimizer:torch.optim.Optimizer,
            loss_fn:torch.nn.Module,
            device: torch.device) -> tuple:
    train_accuracy_hist = [ ]
    test_accuracy_hist = [ ]
    train_loss_hist = [ ]
    test_loss_hist = [ ]
    for t in range(iterations):
        model.train()
        accuracy = 0
        running_loss = 0.0
        for _, data in enumerate(trainloader, 0):
            inputs, labels = data
            x_train, y_train = inputs.to(device), labels.to(device)
            y_pred = model(x_train)
            loss_val = loss_fn(y_pred, y_train)
            running_loss += loss_val.item()
            optimizer.zero_grad()
            loss_val.backward()
            optimizer.step()
            _, predicted = torch.max(y_pred, 1)
            accuracy += (predicted == y_train).sum().item()
        train_accuracy_hist.append(accuracy / len(trainloader.dataset))
        train_loss_hist.append(running_loss / len(trainloader))
        model.eval()
        with torch.inference_mode():
            accuracy = 0
            running_loss = 0.0
            for i, data in enumerate(testloader, 0):
                inputs, labels = data
                x_test, y_test = inputs.to(device), labels.to(device)
                y_pred = model(x_test)
                loss_val = loss_fn(y_pred, y_test)
                running_loss += loss_val.item()
                _, predicted = torch.max(y_pred, 1)
                accuracy += (predicted == y_test).sum().item()

```

```

        test_accuracy_hist.append(accuracy / len(testloader.dataset))
        test_loss_hist.append(running_loss / len(testloader))
        print(f"Epoch {t + 1} / {iterations}, Train Loss:␣
↪{train_loss_hist[-1]}, Test Loss: {test_loss_hist[-1]}, Train Accuracy:␣
↪{train_accuracy_hist[-1]}, Test Accuracy: {test_accuracy_hist[-1]}")
        return train_accuracy_hist, test_accuracy_hist, train_loss_hist,␣
↪test_loss_hist

```

```

[18]: model = NeuralNetwork(Din, 100, K).to(device)
      loss = nn.CrossEntropyLoss()
      optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=reg)
      iterations = 20

```

```

[19]: train_accuracy_hist, test_accuracy_hist, train_loss_hist, test_loss_hist =␣
      ↪train(model, trainloader, testloader, iterations, optimizer, loss, device)

```


```

Epoch 1 / 20, Train Loss: 1.7594276244870684, Test Loss: 1.674093100590447,
Train Accuracy: 0.38396, Test Accuracy: 0.4194
Epoch 2 / 20, Train Loss: 1.6328852262088143, Test Loss: 1.630866586209867,
Train Accuracy: 0.43388, Test Accuracy: 0.4355
Epoch 3 / 20, Train Loss: 1.5735069845856113, Test Loss: 1.592388990968942,
Train Accuracy: 0.45558, Test Accuracy: 0.4492
Epoch 4 / 20, Train Loss: 1.5303320414121533, Test Loss: 1.5951702092021418,
Train Accuracy: 0.46966, Test Accuracy: 0.4508
Epoch 5 / 20, Train Loss: 1.490853445391127, Test Loss: 1.5662361551016664,
Train Accuracy: 0.48536, Test Accuracy: 0.4563
Epoch 6 / 20, Train Loss: 1.4583516220061046, Test Loss: 1.5618095154198595,
Train Accuracy: 0.4966, Test Accuracy: 0.4568
Epoch 7 / 20, Train Loss: 1.426805178209977, Test Loss: 1.547024865119983, Train
Accuracy: 0.50792, Test Accuracy: 0.469
Epoch 8 / 20, Train Loss: 1.403101083985217, Test Loss: 1.5482690997017077,
Train Accuracy: 0.51538, Test Accuracy: 0.4635
Epoch 9 / 20, Train Loss: 1.3720098289250564, Test Loss: 1.5411850149258257,
Train Accuracy: 0.5249, Test Accuracy: 0.4682
Epoch 10 / 20, Train Loss: 1.352620396412723, Test Loss: 1.5606188214244172,
Train Accuracy: 0.53088, Test Accuracy: 0.4726
Epoch 11 / 20, Train Loss: 1.3308184827205392, Test Loss: 1.5377594291592558,
Train Accuracy: 0.53986, Test Accuracy: 0.4712
Epoch 12 / 20, Train Loss: 1.3051582179768149, Test Loss: 1.5428463003505914,
Train Accuracy: 0.54882, Test Accuracy: 0.4692
Epoch 13 / 20, Train Loss: 1.2840144001591, Test Loss: 1.5482081445261313, Train
Accuracy: 0.55636, Test Accuracy: 0.4655
Epoch 14 / 20, Train Loss: 1.2642350733394587, Test Loss: 1.569972787992642,
Train Accuracy: 0.56242, Test Accuracy: 0.4648
Epoch 15 / 20, Train Loss: 1.2450444423000704, Test Loss: 1.541662802330602,
Train Accuracy: 0.57068, Test Accuracy: 0.4714
Epoch 16 / 20, Train Loss: 1.2253786445579236, Test Loss: 1.5717284226189026,
Train Accuracy: 0.57796, Test Accuracy: 0.4673

```

Epoch 17 / 20, Train Loss: 1.2121384872203445, Test Loss: 1.5566994610685891,  
Train Accuracy: 0.58184, Test Accuracy: 0.4682  
Epoch 18 / 20, Train Loss: 1.191947989752105, Test Loss: 1.5680276741996741,  
Train Accuracy: 0.58844, Test Accuracy: 0.4679  
Epoch 19 / 20, Train Loss: 1.175236813356994, Test Loss: 1.5862781468290872,  
Train Accuracy: 0.59512, Test Accuracy: 0.4657  
Epoch 20 / 20, Train Loss: 1.1599616979988279, Test Loss: 1.5828485254662485,  
Train Accuracy: 0.5981, Test Accuracy: 0.469

```
[20]: plt.plot(train_loss_hist, label='Train')  
plt.plot(test_loss_hist, label='Test')  
plt.xlabel('Iteration')  
plt.ylabel('Loss')  
plt.xticks(range(0, iterations, 2))  
plt.title('Loss history')  
plt.legend()  
plt.show()
```




```
[21]: plt.plot(train_accuracy_hist, label='Train')  
plt.plot(test_accuracy_hist, label='Test')
```

```

plt.xlabel('Iteration')
plt.ylabel('Accuracy')
plt.xticks(range(0, iterations, 2))
plt.title('Accuracy history')
plt.legend()
plt.show()

```



```

[22]: def calculate_accuracy(model: nn.Module, dataloader: torch.utils.data.
      ↪ DataLoader) -> float:
      correct = 0
      total = 0
      with torch.no_grad():
          for data in dataloader:
              inputs, labels = data
              x, y = inputs.to(device), labels.to(device)
              outputs = model(x)
              _, predicted = torch.max(outputs, 1)
              total += y.size(0)
              correct += (predicted == y).sum().item()

```

```
return 100 * correct / total
```

```
[23]: train_accuracy = calculate_accuracy(model, trainloader)
      test_accuracy = calculate_accuracy(model, testloader)

      print(f"Train accuracy: {train_accuracy:.2f}%")
      print(f"Test accuracy: {test_accuracy:.2f}%")
```

Train accuracy: 61.56%

Test accuracy: 46.90%

We see that the accuracy is still very low as was in our custom implementation. As at the time of writing, according to [paperswithcode.com](https://paperswithcode.com), the best accuracy on CIFAR-10 is 99.5%. This is achieved by a model called ViT-H/14 which is a vision transformer. Another thing to note is that the model is beginning to overfit after just 3 epochs. This is because the model is too simple and is not able to learn the complex patterns in the data.

```
[ ]: del model, trainloader, testloader, train_accuracy_hist, test_accuracy_hist, \
      ↪ train_loss_hist, test_loss_hist
      gc.collect()
      if torch.cuda.is_available():
          torch.cuda.empty_cache()
```

## 2 LeNet-5

Here we will be implementing LeNet-5 architecture for MNIST dataset.

```
[24]: batch_size = 32
```

```
[25]: trainset_mnist = torchvision.datasets.MNIST(root='./data', train=True,
        ↳download=True, transform=transforms.ToTensor())
trainloader_mnist = torch.utils.data.DataLoader(trainset_mnist,
        ↳batch_size=batch_size, shuffle=True)
testset_mnist = torchvision.datasets.MNIST(root='./data', train=False,
        ↳download=True, transform=transforms.ToTensor())
testloader_mnist = torch.utils.data.DataLoader(testset_mnist,
        ↳batch_size=batch_size, shuffle=False)
classes = tuple(str(i) for i in range(10))
```

### Architecture

```
[26]: class LeNet(nn.Module):
        def __init__(self, input_size, input_channels, output_size):
            super(LeNet, self).__init__()
            self.conv1 = nn.Sequential(
                nn.Conv2d(input_channels, 6, 5),
                nn.ReLU(),
                nn.MaxPool2d(2)
            )
            self.conv2 = nn.Sequential(
                nn.Conv2d(6, 16, 5),
                nn.ReLU(),
                nn.MaxPool2d(2)
            )

            conv_output_size = ((input_size - 4) // 2 - 4) // 2
            self.classifier = nn.Sequential(
                nn.Linear(16 * conv_output_size * conv_output_size, 120),
                nn.ReLU(),
                nn.Linear(120, 84),
                nn.ReLU(),
                nn.Linear(84, output_size)
            )

        def forward(self, x):
            y = self.conv1(x)
            y = self.conv2(y)
            y = y.view(y.size(0), -1)
            y = self.classifier(y)
            return y
```


```
[34]: lenet_model = LeNet(input_size = 28, input_channels = 1, output_size = 10).
      ↪to(device)
      loss = nn.CrossEntropyLoss()
      optimizer = optim.Adam(lenet_model.parameters(), lr=0.001)
      iterations = 5 # Sufficient since MNIST is a simple dataset
```

```
[35]: train_accuracy_hist, test_accuracy_hist, train_loss_hist, test_loss_hist =
      ↪train(lenet_model, trainloader_mnist, testloader_mnist, iterations,
      ↪optimizer, loss, device)
```


```
Epoch 1 / 5, Train Loss: 0.23151897346563638, Test Loss: 0.07052079150938584,
Train Accuracy: 0.92745, Test Accuracy: 0.9773
Epoch 2 / 5, Train Loss: 0.07228330143981923, Test Loss: 0.052754656721598105,
Train Accuracy: 0.9766833333333333, Test Accuracy: 0.9839
Epoch 3 / 5, Train Loss: 0.054350806503665326, Test Loss: 0.05420349845966769,
Train Accuracy: 0.9831666666666666, Test Accuracy: 0.9835
Epoch 4 / 5, Train Loss: 0.04200671799731596, Test Loss: 0.036045840951620965,
Train Accuracy: 0.98675, Test Accuracy: 0.9879
Epoch 5 / 5, Train Loss: 0.035717702910013034, Test Loss: 0.037477730094514015,
Train Accuracy: 0.9888166666666667, Test Accuracy: 0.9882
```

```
[36]: plt.plot(train_accuracy_hist, label='Train Accuracy')
      plt.plot(test_accuracy_hist, label='Test Accuracy')
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
      plt.legend()
      plt.show()
```





```
[37]: plt.plot(train_loss_hist, label='Train Loss')
plt.plot(test_loss_hist, label='Test Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
[38]: train_accuracy = calculate_accuracy(lenet_model, trainloader_mnist)
      test_accuracy = calculate_accuracy(lenet_model, testloader_mnist)

      print(f"Train accuracy: {train_accuracy:.2f}%")
      print(f"Test accuracy: {test_accuracy:.2f}%")
```

Train accuracy: 99.33%

Test accuracy: 98.82%

Observing the plots of loss and accuracy, we can see that the model is performing well. As expected, the train and test losses are decreasing and the train and test accuracies are increasing with each epoch. After 5 epochs, the model was able to achieve a test accuracy of 98.82%. This is easy to achieve as the MNIST dataset is simple and LeNet-5 is a good architecture for this dataset. 5 epochs were used since the model began to overfit after this point.

### 3 Implementing ResNet-18

In this section, we will implement ResNet-18 architecture for classifying the hymenoptera dataset consisting of images of ants and bees. In this first section, we will be finetuning the network where we will be using a pre-trained model and retraining it on the hymenoptera dataset. In the second section, we will be using the network as a feature extractor where we freeze the weights of the network and only train the final classification layer.

#### 3.1 Finetuning the network

```
[59]: resnet_model = torchvision.models.resnet18(weights = 'IMAGENET1K_V1').
      ↪to(device) # These are the default weights
data_folder = './data/hymenoptera_data'
train_transforms = transforms.Compose([transforms.RandomResizedCrop(224),
      ↪transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.
      ↪Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])
trainset_hymenoptera = torchvision.datasets.ImageFolder(root=f'{data_folder}/
      ↪train', transform=train_transforms)
trainloader_hymenoptera = torch.utils.data.DataLoader(trainset_hymenoptera,
      ↪batch_size=batch_size, shuffle=True)
test_transforms = transforms.Compose([transforms.Resize(256), transforms.
      ↪CenterCrop(224), transforms.ToTensor(), transforms.Normalize([0.485, 0.456,
      ↪0.406], [0.229, 0.224, 0.225])])
testset_hymenoptera = torchvision.datasets.ImageFolder(root=f'{data_folder}/
      ↪val', transform=test_transforms)
testloader_hymenoptera = torch.utils.data.DataLoader(testset_hymenoptera,
      ↪batch_size=batch_size, shuffle=False)
classes_hymenoptera = trainset_hymenoptera.classes
```

```
[60]: print(summary(resnet_model, input_size=(batch_size, 3, 224, 224),
      ↪col_names=["input_size", "output_size", "num_params", "trainable"]))
```

```
=====
=====
```

Layer (type:depth-idx)		Input Shape	Output Shape
Param #	Trainable		
=====			
ResNet		[32, 3, 224, 224]	[32, 1000]
--	True		
Conv2d: 1-1		[32, 3, 224, 224]	[32, 64, 112,
112]	9,408	True	
BatchNorm2d: 1-2		[32, 64, 112, 112]	[32, 64, 112,
112]	128	True	
ReLU: 1-3		[32, 64, 112, 112]	[32, 64, 112,
112]	--	--	
MaxPool2d: 1-4		[32, 64, 112, 112]	[32, 64, 56,
56]	--	--	

Sequential: 1-5	[32, 64, 56, 56]	[32, 64, 56,
56]       --	True	
BasicBlock: 2-1	[32, 64, 56, 56]	[32, 64, 56,
56]       --	True	
Conv2d: 3-1	[32, 64, 56, 56]	[32, 64, 56,
56]       36,864	True	
BatchNorm2d: 3-2	[32, 64, 56, 56]	[32, 64, 56,
56]       128	True	
ReLU: 3-3	[32, 64, 56, 56]	[32, 64, 56,
56]       --	--	
Conv2d: 3-4	[32, 64, 56, 56]	[32, 64, 56,
56]       36,864	True	
BatchNorm2d: 3-5	[32, 64, 56, 56]	[32, 64, 56,
56]       128	True	
ReLU: 3-6	[32, 64, 56, 56]	[32, 64, 56,
56]       --	--	
BasicBlock: 2-2	[32, 64, 56, 56]	[32, 64, 56,
56]       --	True	
Conv2d: 3-7	[32, 64, 56, 56]	[32, 64, 56,
56]       36,864	True	
BatchNorm2d: 3-8	[32, 64, 56, 56]	[32, 64, 56,
56]       128	True	
ReLU: 3-9	[32, 64, 56, 56]	[32, 64, 56,
56]       --	--	
Conv2d: 3-10	[32, 64, 56, 56]	[32, 64, 56,
56]       36,864	True	
BatchNorm2d: 3-11	[32, 64, 56, 56]	[32, 64, 56,
56]       128	True	
ReLU: 3-12	[32, 64, 56, 56]	[32, 64, 56,
56]       --	--	
Sequential: 1-6	[32, 64, 56, 56]	[32, 128, 28,
28]       --	True	
BasicBlock: 2-3	[32, 64, 56, 56]	[32, 128, 28,
28]       --	True	
Conv2d: 3-13	[32, 64, 56, 56]	[32, 128, 28,
28]       73,728	True	
BatchNorm2d: 3-14	[32, 128, 28, 28]	[32, 128, 28,
28]       256	True	
ReLU: 3-15	[32, 128, 28, 28]	[32, 128, 28,
28]       --	--	
Conv2d: 3-16	[32, 128, 28, 28]	[32, 128, 28,
28]       147,456	True	
BatchNorm2d: 3-17	[32, 128, 28, 28]	[32, 128, 28,
28]       256	True	
Sequential: 3-18	[32, 64, 56, 56]	[32, 128, 28,
28]       8,448	True	
ReLU: 3-19	[32, 128, 28, 28]	[32, 128, 28,
28]       --	--	

	BasicBlock: 2-4	[32, 128, 28, 28]	[32, 128, 28,
28]	--	True	
	Conv2d: 3-20	[32, 128, 28, 28]	[32, 128, 28,
28]	147,456	True	
	BatchNorm2d: 3-21	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	ReLU: 3-22	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Conv2d: 3-23	[32, 128, 28, 28]	[32, 128, 28,
28]	147,456	True	
	BatchNorm2d: 3-24	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	ReLU: 3-25	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Sequential: 1-7	[32, 128, 28, 28]	[32, 256, 14,
14]	--	True	
	BasicBlock: 2-5	[32, 128, 28, 28]	[32, 256, 14,
14]	--	True	
	Conv2d: 3-26	[32, 128, 28, 28]	[32, 256, 14,
14]	294,912	True	
	BatchNorm2d: 3-27	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-28	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Conv2d: 3-29	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-30	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	Sequential: 3-31	[32, 128, 28, 28]	[32, 256, 14,
14]	33,280	True	
	ReLU: 3-32	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	BasicBlock: 2-6	[32, 256, 14, 14]	[32, 256, 14,
14]	--	True	
	Conv2d: 3-33	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-34	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-35	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Conv2d: 3-36	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-37	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-38	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Sequential: 1-8	[32, 256, 14, 14]	[32, 512, 7,
7]	--	True	

	BasicBlock: 2-7	[32, 256, 14, 14]	[32, 512, 7,
7]	--	True	
	Conv2d: 3-39	[32, 256, 14, 14]	[32, 512, 7,
7]	1,179,648	True	
	BatchNorm2d: 3-40	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	ReLU: 3-41	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	Conv2d: 3-42	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-43	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	Sequential: 3-44	[32, 256, 14, 14]	[32, 512, 7,
7]	132,096	True	
	ReLU: 3-45	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	BasicBlock: 2-8	[32, 512, 7, 7]	[32, 512, 7,
7]	--	True	
	Conv2d: 3-46	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-47	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	ReLU: 3-48	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	Conv2d: 3-49	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-50	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	ReLU: 3-51	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	AdaptiveAvgPool2d: 1-9	[32, 512, 7, 7]	[32, 512, 1,
1]	--	--	
	Linear: 1-10	[32, 512]	[32, 1000]
513,000	True		

```

=====
Total params: 11,689,512
Trainable params: 11,689,512
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 58.05
=====
=====

```

```

Input size (MB): 19.27
Forward/backward pass size (MB): 1271.92
Params size (MB): 46.76
Estimated Total Size (MB): 1337.94
=====
=====

```

```
[61]: resnet_model.fc = nn.Linear(512, len(classes_hymenoptera)).to(device)
```

```
[62]: print(summary(resnet_model, input_size=(batch_size, 3, 224, 224),
    ↪col_names=["input_size", "output_size", "num_params", "trainable"]))
```

```
=====
=====
```

Layer (type:depth-idx)	Input Shape	Output Shape
Param #	Trainable	
ResNet	[32, 3, 224, 224]	[32, 2]
--	True	
Conv2d: 1-1	[32, 3, 224, 224]	[32, 64, 112,
112] 9,408	True	
BatchNorm2d: 1-2	[32, 64, 112, 112]	[32, 64, 112,
112] 128	True	
ReLU: 1-3	[32, 64, 112, 112]	[32, 64, 112,
112] --	--	
MaxPool2d: 1-4	[32, 64, 112, 112]	[32, 64, 56,
56] --	--	
Sequential: 1-5	[32, 64, 56, 56]	[32, 64, 56,
56] --	True	
BasicBlock: 2-1	[32, 64, 56, 56]	[32, 64, 56,
56] --	True	
Conv2d: 3-1	[32, 64, 56, 56]	[32, 64, 56,
56] 36,864	True	
BatchNorm2d: 3-2	[32, 64, 56, 56]	[32, 64, 56,
56] 128	True	
ReLU: 3-3	[32, 64, 56, 56]	[32, 64, 56,
56] --	--	
Conv2d: 3-4	[32, 64, 56, 56]	[32, 64, 56,
56] 36,864	True	
BatchNorm2d: 3-5	[32, 64, 56, 56]	[32, 64, 56,
56] 128	True	
ReLU: 3-6	[32, 64, 56, 56]	[32, 64, 56,
56] --	--	
BasicBlock: 2-2	[32, 64, 56, 56]	[32, 64, 56,
56] --	True	
Conv2d: 3-7	[32, 64, 56, 56]	[32, 64, 56,
56] 36,864	True	
BatchNorm2d: 3-8	[32, 64, 56, 56]	[32, 64, 56,
56] 128	True	
ReLU: 3-9	[32, 64, 56, 56]	[32, 64, 56,
56] --	--	
Conv2d: 3-10	[32, 64, 56, 56]	[32, 64, 56,
56] 36,864	True	
BatchNorm2d: 3-11	[32, 64, 56, 56]	[32, 64, 56,

56]	128	True	
	ReLU: 3-12	[32, 64, 56, 56]	[32, 64, 56,
56]	--	--	
	Sequential: 1-6	[32, 64, 56, 56]	[32, 128, 28,
28]	--	True	
	BasicBlock: 2-3	[32, 64, 56, 56]	[32, 128, 28,
28]	--	True	
	Conv2d: 3-13	[32, 64, 56, 56]	[32, 128, 28,
28]	73,728	True	
	BatchNorm2d: 3-14	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	ReLU: 3-15	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Conv2d: 3-16	[32, 128, 28, 28]	[32, 128, 28,
28]	147,456	True	
	BatchNorm2d: 3-17	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	Sequential: 3-18	[32, 64, 56, 56]	[32, 128, 28,
28]	8,448	True	
	ReLU: 3-19	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	BasicBlock: 2-4	[32, 128, 28, 28]	[32, 128, 28,
28]	--	True	
	Conv2d: 3-20	[32, 128, 28, 28]	[32, 128, 28,
28]	147,456	True	
	BatchNorm2d: 3-21	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	ReLU: 3-22	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Conv2d: 3-23	[32, 128, 28, 28]	[32, 128, 28,
28]	147,456	True	
	BatchNorm2d: 3-24	[32, 128, 28, 28]	[32, 128, 28,
28]	256	True	
	ReLU: 3-25	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Sequential: 1-7	[32, 128, 28, 28]	[32, 256, 14,
14]	--	True	
	BasicBlock: 2-5	[32, 128, 28, 28]	[32, 256, 14,
14]	--	True	
	Conv2d: 3-26	[32, 128, 28, 28]	[32, 256, 14,
14]	294,912	True	
	BatchNorm2d: 3-27	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-28	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Conv2d: 3-29	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-30	[32, 256, 14, 14]	[32, 256, 14,



14]	512	True	
	Sequential: 3-31	[32, 128, 28, 28]	[32, 256, 14,
14]	33,280	True	
	ReLU: 3-32	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	BasicBlock: 2-6	[32, 256, 14, 14]	[32, 256, 14,
14]	--	True	
	Conv2d: 3-33	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-34	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-35	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Conv2d: 3-36	[32, 256, 14, 14]	[32, 256, 14,
14]	589,824	True	
	BatchNorm2d: 3-37	[32, 256, 14, 14]	[32, 256, 14,
14]	512	True	
	ReLU: 3-38	[32, 256, 14, 14]	[32, 256, 14,
14]	--	--	
	Sequential: 1-8	[32, 256, 14, 14]	[32, 512, 7,
7]	--	True	
	BasicBlock: 2-7	[32, 256, 14, 14]	[32, 512, 7,
7]	--	True	
	Conv2d: 3-39	[32, 256, 14, 14]	[32, 512, 7,
7]	1,179,648	True	
	BatchNorm2d: 3-40	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	ReLU: 3-41	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	Conv2d: 3-42	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-43	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	Sequential: 3-44	[32, 256, 14, 14]	[32, 512, 7,
7]	132,096	True	
	ReLU: 3-45	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	BasicBlock: 2-8	[32, 512, 7, 7]	[32, 512, 7,
7]	--	True	
	Conv2d: 3-46	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-47	[32, 512, 7, 7]	[32, 512, 7,
7]	1,024	True	
	ReLU: 3-48	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	Conv2d: 3-49	[32, 512, 7, 7]	[32, 512, 7,
7]	2,359,296	True	
	BatchNorm2d: 3-50	[32, 512, 7, 7]	[32, 512, 7,

```

7]          1,024          True
          ReLU: 3-51      [32, 512, 7, 7]      [32, 512, 7,
7]          --          --
          AdaptiveAvgPool2d: 1-9      [32, 512, 7, 7]      [32, 512, 1,
1]          --          --
          Linear: 1-10      [32, 512]      [32, 2]
1,026          True
=====
=====
Total params: 11,177,538
Trainable params: 11,177,538
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 58.03
=====
=====
Input size (MB): 19.27
Forward/backward pass size (MB): 1271.66
Params size (MB): 44.71
Estimated Total Size (MB): 1335.64
=====
=====

```

```

[63]: loss = nn.CrossEntropyLoss()
      optimizer = optim.SGD(resnet_model.parameters(), lr=0.001)
      iterations = 30

```

```

[64]: train_accuracy_hist, test_accuracy_hist, train_loss_hist, test_loss_hist =
      ↪train(resnet_model, trainloader_hymenoptera, testloader_hymenoptera,
      ↪iterations, optimizer, loss, device)

```

```


Epoch 1 / 30, Train Loss: 0.740321509540081, Test Loss: 0.6342014193534851,
Train Accuracy: 0.5491803278688525, Test Accuracy: 0.6143790849673203
Epoch 2 / 30, Train Loss: 0.6432950645685196, Test Loss: 0.5895971953868866,
Train Accuracy: 0.6311475409836066, Test Accuracy: 0.6928104575163399
Epoch 3 / 30, Train Loss: 0.5841600596904755, Test Loss: 0.5460123479366302,
Train Accuracy: 0.7377049180327869, Test Accuracy: 0.738562091503268
Epoch 4 / 30, Train Loss: 0.5604859106242657, Test Loss: 0.5282542407512665,
Train Accuracy: 0.7540983606557377, Test Accuracy: 0.7712418300653595
Epoch 5 / 30, Train Loss: 0.5212857164442539, Test Loss: 0.49475630521774294,
Train Accuracy: 0.7950819672131147, Test Accuracy: 0.8366013071895425
Epoch 6 / 30, Train Loss: 0.4841243512928486, Test Loss: 0.46025643348693845,
Train Accuracy: 0.8237704918032787, Test Accuracy: 0.8431372549019608
Epoch 7 / 30, Train Loss: 0.44586482644081116, Test Loss: 0.43508976697921753,
Train Accuracy: 0.8524590163934426, Test Accuracy: 0.8627450980392157
Epoch 8 / 30, Train Loss: 0.43780064582824707, Test Loss: 0.41472959518432617,
Train Accuracy: 0.8565573770491803, Test Accuracy: 0.8823529411764706
Epoch 9 / 30, Train Loss: 0.43369603529572487, Test Loss: 0.3919881582260132,
Train Accuracy: 0.8770491803278688, Test Accuracy: 0.8888888888888888

```


Epoch 10 / 30, Train Loss: 0.4077473431825638, Test Loss: 0.3754044145345688,  
 Train Accuracy: 0.860655737704918, Test Accuracy: 0.8823529411764706  
 Epoch 11 / 30, Train Loss: 0.40262167155742645, Test Loss: 0.36203463077545167,  
 Train Accuracy: 0.889344262295082, Test Accuracy: 0.8954248366013072  
 Epoch 12 / 30, Train Loss: 0.35084324330091476, Test Loss: 0.3527695506811142,  
 Train Accuracy: 0.9098360655737705, Test Accuracy: 0.8888888888888888  
 Epoch 13 / 30, Train Loss: 0.3614560030400753, Test Loss: 0.34041716158390045,  
 Train Accuracy: 0.9139344262295082, Test Accuracy: 0.8954248366013072  
 Epoch 14 / 30, Train Loss: 0.34818144142627716, Test Loss: 0.3257443279027939,  
 Train Accuracy: 0.9098360655737705, Test Accuracy: 0.9084967320261438  
 Epoch 15 / 30, Train Loss: 0.38046708330512047, Test Loss: 0.31786433458328245,  
 Train Accuracy: 0.8729508196721312, Test Accuracy: 0.9019607843137255  
 Epoch 16 / 30, Train Loss: 0.320930652320385, Test Loss: 0.3091241180896759,  
 Train Accuracy: 0.9139344262295082, Test Accuracy: 0.9019607843137255  
 Epoch 17 / 30, Train Loss: 0.3131910301744938, Test Loss: 0.3030261009931564,  
 Train Accuracy: 0.9221311475409836, Test Accuracy: 0.9084967320261438  
 Epoch 18 / 30, Train Loss: 0.31882288306951523, Test Loss: 0.2949344992637634,  
 Train Accuracy: 0.9098360655737705, Test Accuracy: 0.9150326797385621  
 Epoch 19 / 30, Train Loss: 0.3195708766579628, Test Loss: 0.29169304966926574,  
 Train Accuracy: 0.9057377049180327, Test Accuracy: 0.9084967320261438  
 Epoch 20 / 30, Train Loss: 0.2938222214579582, Test Loss: 0.28420203030109403,  
 Train Accuracy: 0.9221311475409836, Test Accuracy: 0.9084967320261438  
 Epoch 21 / 30, Train Loss: 0.2851744070649147, Test Loss: 0.27806246280670166,  
 Train Accuracy: 0.9180327868852459, Test Accuracy: 0.9150326797385621  
 Epoch 22 / 30, Train Loss: 0.27146914787590504, Test Loss: 0.2746284455060959,  
 Train Accuracy: 0.9262295081967213, Test Accuracy: 0.9150326797385621  
 Epoch 23 / 30, Train Loss: 0.27277664467692375, Test Loss: 0.26869991421699524,  
 Train Accuracy: 0.9426229508196722, Test Accuracy: 0.9150326797385621  
 Epoch 24 / 30, Train Loss: 0.2761826105415821, Test Loss: 0.26698081791400907,  
 Train Accuracy: 0.9262295081967213, Test Accuracy: 0.9215686274509803  
 Epoch 25 / 30, Train Loss: 0.2681906670331955, Test Loss: 0.2616105377674103,  
 Train Accuracy: 0.9016393442622951, Test Accuracy: 0.9215686274509803  
 Epoch 26 / 30, Train Loss: 0.24914774484932423, Test Loss: 0.2588020324707031,  
 Train Accuracy: 0.9549180327868853, Test Accuracy: 0.9215686274509803  
 Epoch 27 / 30, Train Loss: 0.24066543392837048, Test Loss: 0.25609440803527833,  
 Train Accuracy: 0.9426229508196722, Test Accuracy: 0.9150326797385621  
 Epoch 28 / 30, Train Loss: 0.250135937705636, Test Loss: 0.2535726338624954,  
 Train Accuracy: 0.930327868852459, Test Accuracy: 0.9215686274509803  
 Epoch 29 / 30, Train Loss: 0.23681160807609558, Test Loss: 0.24619961082935332,  
 Train Accuracy: 0.9426229508196722, Test Accuracy: 0.9150326797385621  
 Epoch 30 / 30, Train Loss: 0.2245215941220522, Test Loss: 0.24365059286355972,  
 Train Accuracy: 0.9426229508196722, Test Accuracy: 0.9215686274509803

```
[65]: plt.plot(train_accuracy_hist, label='Train Accuracy')
      plt.plot(test_accuracy_hist, label='Test Accuracy')
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
```

```
plt.legend()  
plt.show()
```



```
[66]: plt.plot(train_loss_hist, label='Train Loss')  
plt.plot(test_loss_hist, label='Test Loss')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



We can observe the model to be performing well as the train and test losses are decreasing and the train and test accuracies are increasing with each epoch. After 30 epochs, the model was able to achieve a test accuracy of 92.16%.

### 3.2 Using ResNet-18 as a feature extractor

```
[67]: resnet_model = torchvision.models.resnet18(weights = 'IMAGENET1K_V1').
      ↪to(device) # We need to reobtain the model as we have modified it previously
resnet_model.fc = nn.Linear(512, len(classes)).to(device)
for param in resnet_model.parameters():
    param.requires_grad = False

for param in resnet_model.fc.parameters():
    param.requires_grad = True
```

```
[68]: loss = nn.CrossEntropyLoss()
optimizer = optim.SGD(resnet_model.parameters(), lr=0.001)
iterations = 30
```

```
[69]: print(summary(resnet_model, input_size=(batch_size, 3, 224, 224),
      ↪col_names=["input_size", "output_size", "num_params", "trainable"]))
```

```
=====
=====
```

Layer (type:depth-idx)		Input Shape	Output Shape
Param #	Trainable		
=====			
ResNet		[32, 3, 224, 224]	[32, 10]
--	Partial		
Conv2d: 1-1		[32, 3, 224, 224]	[32, 64, 112,
112] (9,408)	False		
BatchNorm2d: 1-2		[32, 64, 112, 112]	[32, 64, 112,
112] (128)	False		
ReLU: 1-3		[32, 64, 112, 112]	[32, 64, 112,
112] --	--		
MaxPool2d: 1-4		[32, 64, 112, 112]	[32, 64, 56,
56] --	--		
Sequential: 1-5		[32, 64, 56, 56]	[32, 64, 56,
56] --	False		
BasicBlock: 2-1		[32, 64, 56, 56]	[32, 64, 56,
56] --	False		
Conv2d: 3-1		[32, 64, 56, 56]	[32, 64, 56,
56] (36,864)	False		
BatchNorm2d: 3-2		[32, 64, 56, 56]	[32, 64, 56,
56] (128)	False		
ReLU: 3-3		[32, 64, 56, 56]	[32, 64, 56,
56] --	--		
Conv2d: 3-4		[32, 64, 56, 56]	[32, 64, 56,
56] (36,864)	False		
BatchNorm2d: 3-5		[32, 64, 56, 56]	[32, 64, 56,
56] (128)	False		

	ReLU: 3-6	[32, 64, 56, 56]	[32, 64, 56,
56]	--	--	
	BasicBlock: 2-2	[32, 64, 56, 56]	[32, 64, 56,
56]	--	False	
	Conv2d: 3-7	[32, 64, 56, 56]	[32, 64, 56,
56]	(36,864)	False	
	BatchNorm2d: 3-8	[32, 64, 56, 56]	[32, 64, 56,
56]	(128)	False	
	ReLU: 3-9	[32, 64, 56, 56]	[32, 64, 56,
56]	--	--	
	Conv2d: 3-10	[32, 64, 56, 56]	[32, 64, 56,
56]	(36,864)	False	
	BatchNorm2d: 3-11	[32, 64, 56, 56]	[32, 64, 56,
56]	(128)	False	
	ReLU: 3-12	[32, 64, 56, 56]	[32, 64, 56,
56]	--	--	
	Sequential: 1-6	[32, 64, 56, 56]	[32, 128, 28,
28]	--	False	
	BasicBlock: 2-3	[32, 64, 56, 56]	[32, 128, 28,
28]	--	False	
	Conv2d: 3-13	[32, 64, 56, 56]	[32, 128, 28,
28]	(73,728)	False	
	BatchNorm2d: 3-14	[32, 128, 28, 28]	[32, 128, 28,
28]	(256)	False	
	ReLU: 3-15	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Conv2d: 3-16	[32, 128, 28, 28]	[32, 128, 28,
28]	(147,456)	False	
	BatchNorm2d: 3-17	[32, 128, 28, 28]	[32, 128, 28,
28]	(256)	False	
	Sequential: 3-18	[32, 64, 56, 56]	[32, 128, 28,
28]	(8,448)	False	
	ReLU: 3-19	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	BasicBlock: 2-4	[32, 128, 28, 28]	[32, 128, 28,
28]	--	False	
	Conv2d: 3-20	[32, 128, 28, 28]	[32, 128, 28,
28]	(147,456)	False	
	BatchNorm2d: 3-21	[32, 128, 28, 28]	[32, 128, 28,
28]	(256)	False	
	ReLU: 3-22	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	
	Conv2d: 3-23	[32, 128, 28, 28]	[32, 128, 28,
28]	(147,456)	False	
	BatchNorm2d: 3-24	[32, 128, 28, 28]	[32, 128, 28,
28]	(256)	False	
	ReLU: 3-25	[32, 128, 28, 28]	[32, 128, 28,
28]	--	--	

Sequential: 1-7	[32, 128, 28, 28]	[32, 256, 14,
14] --	False	
BasicBlock: 2-5	[32, 128, 28, 28]	[32, 256, 14,
14] --	False	
Conv2d: 3-26	[32, 128, 28, 28]	[32, 256, 14,
14] (294,912)	False	
BatchNorm2d: 3-27	[32, 256, 14, 14]	[32, 256, 14,
14] (512)	False	
ReLU: 3-28	[32, 256, 14, 14]	[32, 256, 14,
14] --	--	
Conv2d: 3-29	[32, 256, 14, 14]	[32, 256, 14,
14] (589,824)	False	
BatchNorm2d: 3-30	[32, 256, 14, 14]	[32, 256, 14,
14] (512)	False	
Sequential: 3-31	[32, 128, 28, 28]	[32, 256, 14,
14] (33,280)	False	
ReLU: 3-32	[32, 256, 14, 14]	[32, 256, 14,
14] --	--	
BasicBlock: 2-6	[32, 256, 14, 14]	[32, 256, 14,
14] --	False	
Conv2d: 3-33	[32, 256, 14, 14]	[32, 256, 14,
14] (589,824)	False	
BatchNorm2d: 3-34	[32, 256, 14, 14]	[32, 256, 14,
14] (512)	False	
ReLU: 3-35	[32, 256, 14, 14]	[32, 256, 14,
14] --	--	
Conv2d: 3-36	[32, 256, 14, 14]	[32, 256, 14,
14] (589,824)	False	
BatchNorm2d: 3-37	[32, 256, 14, 14]	[32, 256, 14,
14] (512)	False	
ReLU: 3-38	[32, 256, 14, 14]	[32, 256, 14,
14] --	--	
Sequential: 1-8	[32, 256, 14, 14]	[32, 512, 7,
7] --	False	
BasicBlock: 2-7	[32, 256, 14, 14]	[32, 512, 7,
7] --	False	
Conv2d: 3-39	[32, 256, 14, 14]	[32, 512, 7,
7] (1,179,648)	False	
BatchNorm2d: 3-40	[32, 512, 7, 7]	[32, 512, 7,
7] (1,024)	False	
ReLU: 3-41	[32, 512, 7, 7]	[32, 512, 7,
7] --	--	
Conv2d: 3-42	[32, 512, 7, 7]	[32, 512, 7,
7] (2,359,296)	False	
BatchNorm2d: 3-43	[32, 512, 7, 7]	[32, 512, 7,
7] (1,024)	False	
Sequential: 3-44	[32, 256, 14, 14]	[32, 512, 7,
7] (132,096)	False	



	ReLU: 3-45	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	BasicBlock: 2-8	[32, 512, 7, 7]	[32, 512, 7,
7]	--	False	
	Conv2d: 3-46	[32, 512, 7, 7]	[32, 512, 7,
7]	(2,359,296)	False	
	BatchNorm2d: 3-47	[32, 512, 7, 7]	[32, 512, 7,
7]	(1,024)	False	
	ReLU: 3-48	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	Conv2d: 3-49	[32, 512, 7, 7]	[32, 512, 7,
7]	(2,359,296)	False	
	BatchNorm2d: 3-50	[32, 512, 7, 7]	[32, 512, 7,
7]	(1,024)	False	
	ReLU: 3-51	[32, 512, 7, 7]	[32, 512, 7,
7]	--	--	
	AdaptiveAvgPool2d: 1-9	[32, 512, 7, 7]	[32, 512, 1,
1]	--	--	
	Linear: 1-10	[32, 512]	[32, 10]
5,130	True		

```

=====
Total params: 11,181,642
Trainable params: 5,130
Non-trainable params: 11,176,512
Total mult-adds (Units.GIGABYTES): 58.03
=====
Input size (MB): 19.27
Forward/backward pass size (MB): 1271.66
Params size (MB): 44.73
Estimated Total Size (MB): 1335.66
=====
=====

```

```

[70]: train_accuracy_hist, test_accuracy_hist, train_loss_hist, test_loss_hist =
      ↪train(resnet_model, trainloader_hymenoptera, testloader_hymenoptera,
      ↪iterations, optimizer, loss, device)

```

```


Epoch 1 / 30, Train Loss: 2.169361427426338, Test Loss: 1.623906660079956, Train
Accuracy: 0.18032786885245902, Test Accuracy: 0.39869281045751637
Epoch 2 / 30, Train Loss: 1.4038525372743607, Test Loss: 1.1809960842132567,
Train Accuracy: 0.4713114754098361, Test Accuracy: 0.5032679738562091
Epoch 3 / 30, Train Loss: 1.0933245494961739, Test Loss: 0.982301938533783,
Train Accuracy: 0.5204918032786885, Test Accuracy: 0.5359477124183006
Epoch 4 / 30, Train Loss: 0.9442420229315758, Test Loss: 0.8642782926559448,
Train Accuracy: 0.5532786885245902, Test Accuracy: 0.5555555555555556
Epoch 5 / 30, Train Loss: 0.8431264758110046, Test Loss: 0.785391116142273,

```


Train Accuracy: 0.6065573770491803, Test Accuracy: 0.6339869281045751  
Epoch 6 / 30, Train Loss: 0.8003100976347923, Test Loss: 0.7292196869850158,  
Train Accuracy: 0.6229508196721312, Test Accuracy: 0.673202614379085  
Epoch 7 / 30, Train Loss: 0.7290625125169754, Test Loss: 0.683362352848053,  
Train Accuracy: 0.6926229508196722, Test Accuracy: 0.7124183006535948  
Epoch 8 / 30, Train Loss: 0.6734580993652344, Test Loss: 0.6364186406135559,  
Train Accuracy: 0.7336065573770492, Test Accuracy: 0.7516339869281046  
Epoch 9 / 30, Train Loss: 0.6541433706879616, Test Loss: 0.6063030123710632,  
Train Accuracy: 0.7336065573770492, Test Accuracy: 0.7450980392156863  
Epoch 10 / 30, Train Loss: 0.6121368557214737, Test Loss: 0.5750733911991119,  
Train Accuracy: 0.7827868852459017, Test Accuracy: 0.7908496732026143  
Epoch 11 / 30, Train Loss: 0.5844235718250275, Test Loss: 0.5531553149223327,  
Train Accuracy: 0.7786885245901639, Test Accuracy: 0.8104575163398693  
Epoch 12 / 30, Train Loss: 0.5474961921572685, Test Loss: 0.5237249076366425,  
Train Accuracy: 0.8114754098360656, Test Accuracy: 0.8104575163398693  
Epoch 13 / 30, Train Loss: 0.5326035730540752, Test Loss: 0.502429085969925,  
Train Accuracy: 0.8442622950819673, Test Accuracy: 0.8104575163398693  
Epoch 14 / 30, Train Loss: 0.53502157330513, Test Loss: 0.48578929901123047,  
Train Accuracy: 0.7950819672131147, Test Accuracy: 0.8496732026143791  
Epoch 15 / 30, Train Loss: 0.4917401447892189, Test Loss: 0.471936309337616,  
Train Accuracy: 0.8442622950819673, Test Accuracy: 0.8627450980392157  
Epoch 16 / 30, Train Loss: 0.4841150566935539, Test Loss: 0.4520291805267334,  
Train Accuracy: 0.8114754098360656, Test Accuracy: 0.869281045751634  
Epoch 17 / 30, Train Loss: 0.4584265500307083, Test Loss: 0.43817468285560607,  
Train Accuracy: 0.8483606557377049, Test Accuracy: 0.8823529411764706  
Epoch 18 / 30, Train Loss: 0.46567703410983086, Test Loss: 0.43019127249717715,  
Train Accuracy: 0.860655737704918, Test Accuracy: 0.8954248366013072  
Epoch 19 / 30, Train Loss: 0.44865116477012634, Test Loss: 0.4137004196643829,  
Train Accuracy: 0.8729508196721312, Test Accuracy: 0.8954248366013072  
Epoch 20 / 30, Train Loss: 0.4460241571068764, Test Loss: 0.4021461963653564,  
Train Accuracy: 0.8524590163934426, Test Accuracy: 0.8954248366013072  
Epoch 21 / 30, Train Loss: 0.44250325486063957, Test Loss: 0.3907356262207031,  
Train Accuracy: 0.860655737704918, Test Accuracy: 0.9019607843137255  
Epoch 22 / 30, Train Loss: 0.4419235624372959, Test Loss: 0.38154500126838686,  
Train Accuracy: 0.8319672131147541, Test Accuracy: 0.9150326797385621  
Epoch 23 / 30, Train Loss: 0.4109177030622959, Test Loss: 0.3718034625053406,  
Train Accuracy: 0.8647540983606558, Test Accuracy: 0.9150326797385621  
Epoch 24 / 30, Train Loss: 0.41230393573641777, Test Loss: 0.364398592710495,  
Train Accuracy: 0.889344262295082, Test Accuracy: 0.9215686274509803  
Epoch 25 / 30, Train Loss: 0.4102812893688679, Test Loss: 0.3571122646331787,  
Train Accuracy: 0.8852459016393442, Test Accuracy: 0.9215686274509803  
Epoch 26 / 30, Train Loss: 0.37579767778515816, Test Loss: 0.3505669295787811,  
Train Accuracy: 0.9016393442622951, Test Accuracy: 0.9215686274509803  
Epoch 27 / 30, Train Loss: 0.37975001707673073, Test Loss: 0.3444370269775391,  
Train Accuracy: 0.9016393442622951, Test Accuracy: 0.9215686274509803  
Epoch 28 / 30, Train Loss: 0.3856111168861389, Test Loss: 0.3392278730869293,  
Train Accuracy: 0.8770491803278688, Test Accuracy: 0.9215686274509803  
Epoch 29 / 30, Train Loss: 0.36211057752370834, Test Loss: 0.33200012147426605,

Train Accuracy: 0.8975409836065574, Test Accuracy: 0.9215686274509803  
Epoch 30 / 30, Train Loss: 0.37362760677933693, Test Loss: 0.3237911552190781,  
Train Accuracy: 0.8647540983606558, Test Accuracy: 0.9215686274509803

```
[71]: plt.plot(train_accuracy_hist, label='Train Accuracy')  
plt.plot(test_accuracy_hist, label='Test Accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()
```



```
[72]: plt.plot(train_loss_hist, label='Train Loss')  
plt.plot(test_loss_hist, label='Test Loss')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



When using the model as a feature extractor, the performance is still very good. After 30 epochs, the model was able to achieve a test accuracy of 92.16% which matches the performance when it was finetuned.